

SPEECH MODELS: MEL FREQUENCY CEPSTRAL COEFFICIENTS (MFCC)

Michele Rossi

rossi@dei.unipd.it

Dept. of Information Engineering
University of Padova, IT



Outline

- MFCCs: what are they for?
- The 6 steps to compute them (intuition)
- The Mel scale for human perceived frequency
- The 6 steps in greater detail

MFCC – what are they for?

- First step in any speech recognition system is
 - To **extract features**: should be good for identifying the linguistic content (neglecting background noise, emotion, etc.)
- Sounds generated by humans
 - Are filtered by the shape of the vocal tract, including tongue, teeth, etc.
 - This shape determines which sound comes out
 - If we determine this shape accurately, this should give a good representation of the **phoneme** that is being produced
- Shape of the vocal tract
 - Manifest itself in the envelope of the *short-term power spectrum*
 - The job of MFCCs is to **accurately represent this envelope** [1]

[1] S. Davis, P. Mermelstein, “Comparison of parametric representations for monosyllabic word recognition in continuously spoken sequences,” *IEEE Transactions on Acoustic, Speech and Signal Processing*, Vol. 28, No. 4, Aug. 1980.

Their computation involves 6 steps

1. **Frame** the signal into short frames
2. For each frame: calculate its **power spectrum**
3. Apply the **Mel filterbank** to the power spectrum
 - sum the energy in each filter
4. Take the **logarithm** of all filterbank energies
5. Take the **DCT** of the log filterbank energies
6. **Keep DCT** coefficients 2-13, discard the others

Step 1 - framing

- An audio signal is constantly changing (non-stationary)
- It is difficult to deal with non-stationary data
- To simplify things
 - We assume that on short time scales the signal does not change much
 - “Does not change” means it is **statistically stationary**
- This is why we frame the signal into **20-40ms time windows**
 - Standard value is 25ms
 - Shorter values: not enough samples to get a reliable spectral estimate
 - Larger values: the signal (statistics) changes too much in the frame

Step 2 – power spectrum

- Compute the power spectrum of a frame
 - This is motivated by the human cochlea (organ in the ear)
 - Which vibrates at different spots depending on the frequency of the incoming sound waves
 - Depending on the location of the cochlea that vibrates, different nerves fire informing the brain that certain frequencies are present
- The periodogram estimate of the power spectrum
 - Does a similar job
 - Informing us which frequencies are present in the current frame

Step 3 – filterbank

- The periodogram estimate of the power spectrum
 - Still contains a [lot of information that is not required by Autonomous Speech Recognition \(ASR\)](#) systems
 - In fact, the cochlea cannot discern the difference between two closely spaced frequencies
 - This effect becomes more apparent as frequencies increase
- Hence:
 - We take clumps of periodogram bins and sum them up to get an *idea of how much energy there exists in each region*
 - This is performed using a [Mel filterbank](#)
 - First Mel filter is very narrow: to get an idea of how much energy there is around 0 Hertz
 - Filters get wider as frequencies get higher: we become less concerned about variations (Mel scale tells us exactly how to space filterbanks)

Step 4 – logarithm


- Once we have the filterbank energies
 - We take their logarithm
 - This is also motivated by human hearing: we do not hear loudness on a linear scale
 - Generally, to double the perceived volume of a sound we need to put 8 times as much energy into it
 - This means that large variations in energy *may not sound that different* if the sound is loud
- Hence
 - This compression operation (logarithm) makes our features match more closely to what humans actually hear

Steps 5 and 6 – DCT

- Discrete Cosine Transform (DCT)
 - There are two main reasons for applying it
 - **First:** filterbanks are quite overlapping (meaning that filterbank energies are quite correlated with each other)
 - DCT decorrelated filterbank energies
 - This means that **diagonal covariance matrices** can be used to model the features with GMM and/or GMM/HMM systems (very desirable)
 - Also (energy compaction property of DCT) we obtain a compact representation in the first DCT coefficients
 - **Second:** discard some information for improved performance
 - Only 12 of the 26 DCT coefficients are kept
 - Higher DCT coefficients represent fast changes in the filterbank energies
 - These fast changes degrade ASR performance
 - A (small) improvement is achieved by dropping them

The Mel scale - $M(f)$


- Relates the *human perceived frequency* (or pitch) of a pure tone to its *actual measured frequency*
- Humans are much better at discerning small changes in pitch at low frequencies than they are at high frequencies
- Incorporating it makes ASR match more closely what humans hear



STEPS 1-6 IN MORE DETAIL

Step 1 - framing

- Frame the audio signal into [20,40] ms frames (standard value is 25 ms)
- For a 20.480 kHz signal, with 25 ms we have
 - 20,480 samples per second
 - $0.025 * 20,480 = 512$ samples per frame



- Frame step is usually, e.g., 10 ms (160 samples)
 - This allows for some overlapping between subsequent frames
 - First frame (of 512 samples) starts at sample 0
 - Second frame (still of 512 samples) start at sample 160, etc.
 - Keep framing until the end of the speech signal
 - If the signal does not divide into an even number of frames, pad it with zeros so it does
- Notation
 - $s(n)$ is the audio signal (time domain signal)
 - $s_i(n)$ is frame i with n ranging from 0 to 511 in our example

Step 2 – power spectrum (1/4)

- For each frame $s_i(n)$, $n = 0, 1, \dots, N - 1$ ($N = 512$)
 - Calculate the complex DFT (Discrete Fourier Transform)


$$S_i(k) = \sum_{n=0}^{N-1} s_i(n)w(n)e^{-j2\pi kn/N}, \quad k = 0, 1, \dots, N - 1$$

- **DFT**
 - Converts a time sequence of N equally spaced samples into an N -long complex sequence, which is a complex-valued function of frequency
 - The DFT (with $w(n)=1$) completely describes the discrete Fourier transform onto an **N -periodic time sequence**
 - When applying DFT, **we are implicitly applying it to an infinitely repeating (periodic) signal**. However, if this is not the case, e.g., first and last samples of $s_i(n)$ do not match, this is *interpreted as a discontinuity* in the signal and *generates a lot of high frequency response in the transformed signal*, that we do not want → use of “windowing” **$w(n)$**

Step 2 – power spectrum (2/4)


- Often used: Hamming window - $w(n)$

$$w(n) = 0.54 - 0.46 \cos \left(\frac{2\pi n}{N-1} \right), \quad n = 0, \dots, N-1$$



Step 2 – power spectrum (3/4)

- Hamming window – transformation (with $N=20$)



- The **center of the main lobe** of a smoothing window **occurs at each frequency component of the input signal** (of course, freq. plot is symmetric around zero)
- Center lobe width (across negative & positive freqs.) is $4\Omega_N = 4(2\pi/N)$
- Ideally, only the main lobe should be emphasized while the side lobes should be zeroed (this prevents energy leakage across subsequent frequencies)

Step 2 – power spectrum (4/4)

- Note

- Among the $N=512$ complex DCT values
- Only the first $1+N/2$ values are significant, the remaining ones are complex conjugates of the first $1+N/2$ values (this is how DCT works)
- Hence, the first $1+N/2=257$ elements *are kept*, while the remaining ones ($N/2-1$ values) *are discarded*
- Compute the *periodogram estimate of the power spectrum*:


$$P_i(k) = \frac{|S_i(k)|^2}{N}, \quad k = 0, \dots, N/2$$

Remember that: $k=0$ is the DC component (*frequency $f=0$*). Then, with DCT the step size is F_s/N (F_s is the *sampling frequency*). Hence, $k=1$ corresponds to frequency $f_1=F_s/N$ and element $k=N/2$ corresponds to frequency $f_{N/2}=F_s/2$ (this is the, so called, *Nyquist critical frequency*, i.e., the highest that can be represented by sampling at frequency F_s)


Step 3 – filterbanks (1/3)

- Compute Mel-spaced filterbanks (through Mel's equation)
 - They amounts to 20-40 triangular filters
 - Each filter (in our case) has 257 values (to match the output of step 2) and it is multiplied by the entire power spectrum from step 2


full filterbank (left)
full power spectrum (right)



filter 8 from filterbank (left)
output when this filter is applied (right)



filter 20 from filterbank (left)
output when this filter is applied (right)



Step 3 – filterbanks (2/3)

- Obtain N_{fb} Mel-spaced triangular filters
 - Pick the number of filters (usually $N_{fb}=26$)
 - Set the maximum frequency as $F_{max}=F_s/2$
 - Set the minimum frequency as, e.g., $F_{min}=300$ Hz (user defined)
 - Compute N_{fb} center frequencies $f(1), f(2), \dots, f(N_{fb})$
 - Linearly spaced in Mel's domain
 - Non-linearly spaced in frequency domain [Hz]
- Each filter m
 - Is centered at frequency $f(m)$
 - Is equal to 1 for $f(m)$ and decreasing linearly otherwise
 - Is zero at $f(m-1), f(m+1)$ and outside $[f(m-1), f(m+1)]$
 - N_{fb} filters require $N_{fb}+2$ thresholds



Step 3 – filterbanks (3/3)

- Example

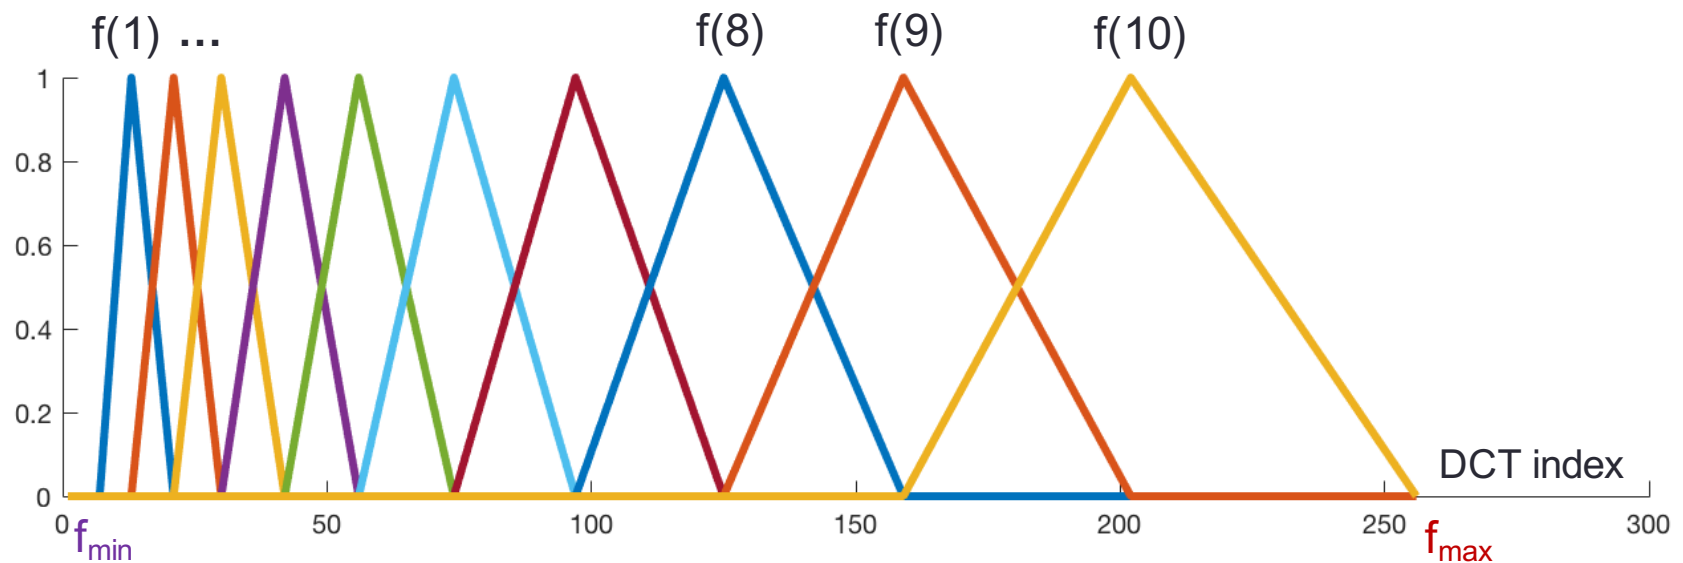
- $F_s = 20.480$ kHz, $N=512$, $N_{fb}=10$, $f_{max}=F_s/2=10,240$, $f_{min}=300$ Hz

- Frequencies [Hz]

$300=f_{min}$ $543=f(1)$ $845=f(2)$ $1,220=f(3)$ $1,687=f(4)$ $2,267=f(5)$ $2,988=f(6)$
 $3,883=f(7)$ $4,997=f(8)$ $6,381=f(9)$ $8,102=f(10)$ $10,240=f_{max}$

- Corresponding DCT indices

[7 13 21 30 42 56 74 97 125 159 202 256]



Step 4 – logarithms


- We now have N_{fb} Mel-spaced triangular filters
- For each filter $m=1,2,\dots, N_{fb}$
 1. Multiply filter m by the full power spectrum from step 2 (the resulting energy trace will be mostly 0 (besides around $f(m)$, the filter acts as a mask, centered at $f(m)$)
 2. Add up the non-zero coefficients in this trace $\rightarrow E_m$
 3. Take the logarithm: $\log(E_m)$
- This gives us N_{fb} real values:
- $\log(E_1), \log(E_2), \dots, \log(E_M)$, with $M= N_{fb}$
 - One for each Mel's frequency band
 - They tell how much energy there is within each band

Steps 5 and 6

- Step 5 – for each frame
 - Take the DCT (Discrete Cosine Transform) of the vector containing the logarithms of the energy within each band
 - The result of the DCT is a vector of N_{fb} cepstral coefficients
- Step 6 – for each frame
 - For ASR, the cepstral coefficients no. 2,3, ...,13 are kept
 - The remaining ones are discarded
 - Note that: the first DCT coefficient is the sum of all the log-energies computed at the previous step (by the very def. of DCT) – thus, it is an overall measure of signal loudness and is not very informative - it is often discarded for *speech recognition* or *speaker id applications* where the system has to be robust to loudness variations
- Final result of this processing is
 - 12 cepstral coefficients for each frame


Some visual insight (1/2)

$F_s = 8$ kHz time signal



Signal in the Time Domain

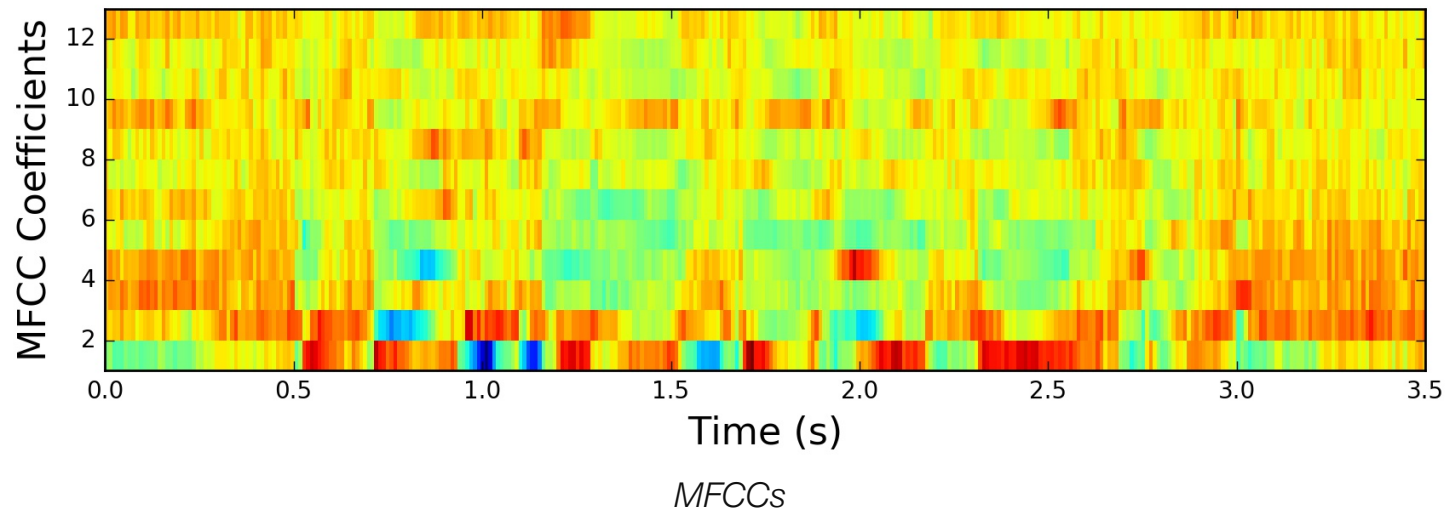
After applying the filterbank to the power spectrum (periodogram):



Spectrogram of the Signal

Some visual insight (2/2)

Mel's Frequency Cepstral Coefficients (MFCCs):



Phyton library to extract MFCCs

https://github.com/jameslyons/python_speech_features

APPENDIX A

Matlab code to compute filterbanks

Filterbanks Matlab code (1/2)

```
Fs=20480;           % sampling rate
nDFT=512;           % number of DFT samples
Fnyq=Fs/2;          % Nyquist frequency

Nfb=10;             % number of filterbanks to generate
Fmin=300;           % min frequency of filterbanks [Hz]
Fmax=Fnyq;          % max frequency of filterbanks [Hz]

FminMel=1125*log(1+(Fmin/700)); % min Mel's frequency
FmaxMel=1125*log(1+(Fmax/700)); % Max Mel's frequency

Nthresh=Nfb+2;      % number of required frequency thresholds

step=(FmaxMel-FminMel)/(Nthresh-1); % step size (linearly partition Mel's space)

% allocate memory for frequency thresholds
melvec=zeros(1,Nthresh);
freqvec=zeros(1,Nthresh);
f=zeros(1,Nthresh);

% for each frequency threshold do
for i = 0:(Nthresh-1)
    melvec(i+1) = FminMel+i*step; % assign Mel thresholds (linearly)
    freqvec(i+1) = 700*(exp(melvec(i+1)/1125)-1); % compute corresponding frequencies (Mel's transformation)
    f(i+1) = floor((nDFT+1)*freqvec(i+1)/Fs); % express frequency thresholds in terms of DFT indices
end
```

Filterbanks Matlab code (2/2)

```
H=zeros(Nfb,Nthresh); % filterbank matrix, here we store the  $N_{fb}$  filters
```

```
% for each filter m do (m is the filterbank index)
```

```
for m = 2:(Nfb+1)
```

```
    for k=1:f(Nthresh)
```

```
        if ((k>=f(m-1)) && (k<=f(m))) H(m-1,k) = (k-f(m-1))/(f(m)-f(m-1)); % increase
```

```
        elseif ((k>=f(m)) && (k<=f(m+1))) H(m-1,k) = (f(m+1)-k)/(f(m+1)-f(m)); % decrease
```

```
        else H(m-1,k) = 0;
```

```
    end
```

```
end
```

```
end
```

```
clf
```

```
figure(1);
```

```
hold on
```

```
for m = 1:Nfb
```

```
    % read Filterbank matrix by row, each row contains a filter
```

```
    plot(H(m,:), 'LineWidth', 3)
```

```
end
```

SPEECH MODELS: MEL FREQUENCY CEPSTRAL COEFFICIENTS (MFCC)

Michele Rossi

rossi@dei.unipd.it

Dept. of Information Engineering
University of Padova, IT

